

PROFILE

Data Engineer with hands on experience building end to end data pipelines using Azure Databricks, Delta Lake, and Medallion Architecture. Skilled in Python, SQL, PySpark, Airflow, Kafka, Microsoft Fabric, and Snowflake. Experienced in ETL/ELT workflows, workflow orchestration, and real time streaming. Proven ability to design structured pipelines and optimize transformations for reliability and performance. Open to opportunities in Data Engineering.

SKILLS

Data Engineering – (Basic Hands on) PySpark, Airflow, Kafka, Azure Databricks, Microsoft Fabric, Snowflake, Delta Lake, Medallion Architecture

Programming – Python, SQL

Data Tools – Pandas, NumPy

Databases – MySQL, Data Warehousing (Snowflake)

Other Interests – Cloud Engineering, Machine Learning, Generative AI.

WORK EXPERIENCE

Evamp & Saanga – Islamabad, Pakistan

Data Engineering Intern | Jul 2025 – Sep 2025

- Strengthened Python and SQL skills for **data analysis, joins, subqueries, and window functions**.
- Implemented **PySpark transformations and ETL pipelines**, applying optimization techniques for distributed processing.
- Built and scheduled **Airflow DAGs** to orchestrate workflows and manage dependencies.
- Worked with **Apache Kafka** to stream data in real-time using producers and consumers.
- Designed **Delta Lake ETL pipelines** and practiced the **Medallion Architecture** on **Azure Databricks**.
- Gained exposure to **Microsoft Fabric** for unified analytics and integration.
- Concluded with **Snowflake** for data warehousing and schema design.

PROJECTS

End-to-End Data Pipeline on Azure Databricks

- Designed an **ETL pipeline** on **Azure Databricks**.
- Implemented **Medallion Architecture (Bronze, Silver, Gold layers)** with **Delta Lake** for structured data processing.
- Ingested raw data into **Bronze**, cleaned and transformed into **Silver**, and delivered curated datasets in **Gold**.
- Built **Power BI dashboards** on top of the Gold layer for reporting and insights.
- Optimized transformations, reducing execution time and improving reliability.

ETL with Microsoft Fabric

- Implemented a data engineering workflow on Microsoft Fabric, integrating data from multiple sources.
- Used Data Factory and Lakehouse features for ingestion and transformation.
- Structured data into Fabric Lakehouse tables, enabling unified access across analytics tools.
- Integrated with Power BI to deliver reports and dashboards on curated data.

Real-Time Data Streaming with Apache Kafka & Airflow

- Ingested data from a **public API** and extracted only the **relevant fields**.
- Implemented a **Kafka producer** (via Kafkaesque) to publish transformed data into topics.
- Orchestrated the pipeline using **Apache Airflow DAGs**, with a **1-minute timer** configured when tasks were manually triggered.
- Demonstrated real-time ingestion and orchestration, showcasing integration of APIs, Kafka, and Airflow.

ACCOMPLISHMENTS

- **Gold Medal (2022, 2023)** – Academic Excellence

EDUCATION

Liverpool John Moore University Intl Program (BIC) – BS Software Engineering (2022 – 2025)

Beaconhouse School System – A Level (2019 – 2021)

CERTIFICATIONS

- Python for Data Science, AI & Development – Coursera
- Data Analysis with Python – Coursera
- Databases and SQL for Data Science with Python – Coursera
- Prepare, Clean, Transform, and Load Data using Power BI – Coursera